

Linear Regression with Multiple Variables

Introduction

- Multiple variables = multiple features
- In the previous session we had
 - X = house size, use this to predict
 - y = house price
- If we consider more variables (such as number of bedrooms, number floors, age of the house) then x_1, x_2, x_3, x_4 are the four features
 - x_1 - size (feet squared)
 - x_2 - Number of bedrooms
 - x_3 - Number of floors
 - x_4 - Age of house (years)
 - y is the output variable (price)

More notations

- **n** - number of features ($n = 4$)
- **m** - number of examples (i.e. number of rows in a table)
- **\mathbf{x}^i** - vector of the input for an example (so a vector of the four parameters for the i^{th} input example)
 - i is an index into the training set
 - So
 - \mathbf{x} is an n -dimensional feature vector
 - \mathbf{x}^3 is, for example, the 3rd house, and contains the four features associated with that house
- **x_j^i** - The value of feature j in the i th training example
 - So
 - x_2^3 is, for example, the number of bedrooms in the third house

With Multiple Features

- What is the form of our **hypothesis**?
- Previously our hypothesis took the form;
 - $h_{\theta}(x) = \theta_0 + \theta_1 x$
 - Here we have two parameters (theta 1 and theta 2) determined by our cost function
 - One variable x
- Now we have multiple features
 - $h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3 + \theta_4 x_4$
- For example
$$h_{\theta}(x) = 80 + 0.1x_1 + 0.01x_2 + 3x_3 - 2x_4$$
 - An example of a hypothesis which is trying to predict the price of a house
 - Parameters are still determined through a **cost function**

- For convenience of notation, $x_0 = 1$ For every example i you have an additional 0th feature for each example
- So now your **feature vector** is $n + 1$ dimensional feature vector indexed from 0
 - This is a column vector called x
 - Each example has a column vector associated with it
 - So let's say we have a new example called " X "
- **Parameters** are also in a 0 indexed $n+1$ dimensional vector
 - This is also a column vector called θ
 - This vector is the same for each example

- Considering this, hypothesis can be written
 - $h_{\theta}(x) = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3 + \theta_4 x_4$
- If we do
 - $h_{\theta}(x) = \theta^T X$
 - θ^T is an $[1 \times n+1]$ matrix
 - In other words, because θ is a column vector, the transposition operation transforms it into a row vector
 - So before
 - θ was a matrix $[n+1 \times 1]$
 - Now
 - θ^T is a matrix $[1 \times n+1]$

- Which means the inner dimensions of θ^T and X match, so they can be multiplied together as
 - $[1 \times n+1] * [n+1 \times 1]$
 - $= h_{\theta}(x)$
 - So, in other words, the transpose of our parameter vector * an input example X gives you a predicted hypothesis which is $[1 \times 1]$ dimensions (i.e. a single value)
- This is an example of multivariate linear regression

Gradient Descent for Multiple Variables

- Fitting parameters for the hypothesis with gradient descent
 - Parameters are θ_0 to θ_n
 - Instead of thinking about this as **n separate values**, think about the parameters as a **single vector (θ)**
 - Where θ is **n+1** dimensional
- The cost function is

$$J(\theta_0, \theta_1, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

- Similarly, instead of thinking of J as a function of the **$n+1$ numbers**, $J()$ is just a function of the parameter vector

- $J(\theta)$

Repeat {

- **Gradient descent** $\rightarrow \theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \dots, \theta_n)$
} (simultaneously update for every $j = 0, \dots, n$)

- Once again, this is
 - $\theta_j = \theta_j - \text{learning rate } (\alpha) \text{ times the partial derivative of } J(\theta)$
with respect to $\theta_{J(\dots)}$
 - We do this through a **simultaneous update** of every θ_j value
- Implementing this algorithm
 - When $n = 1$

Repeat {

$$\theta_0 := \theta_0 - \alpha \underbrace{\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})}_{\frac{\partial}{\partial \theta_0} J(\theta)}$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x^{(i)}$$

(simultaneously update θ_0, θ_1) }

- Above, we have slightly different update rules for θ_0 and θ_1
 - Actually they're the same, except the end has a previously undefined $\mathbf{x}_0^{(i)}$ as $\mathbf{1}$, so wasn't shown
- We now have an almost identical rule for multivariate gradient descent

New algorithm ($n \geq 1$):

Repeat {

$$\downarrow -\frac{2}{2\theta_j} J(\theta)$$

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

(simultaneously update θ_j for
 $j = 0, \dots, n$) }

The interpretation

- We're doing this for each j (o until n) as a simultaneous update (like when $n = 1$)
- So, we re-set θ_j to
 - θ_j minus the learning rate (α) times the partial derivative of the θ vector with respect to θ_j
 - In non-calculus words, this means that we do
 - Learning rate
 - Times $1/m$ (makes the maths easier)
 - Times the sum of
 - The hypothesis taking in the variable vector, minus the actual value, times the j -th value in that variable vector for EACH example

- It's important to remember that

$$\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} = \frac{\partial}{\partial \theta_j} J(\theta)$$

- These algorithm are highly similar

Next

- Feature scaling (normalization)
- How to choose a learning rate
- How to check if GD is working?
- Features and polynomial regression